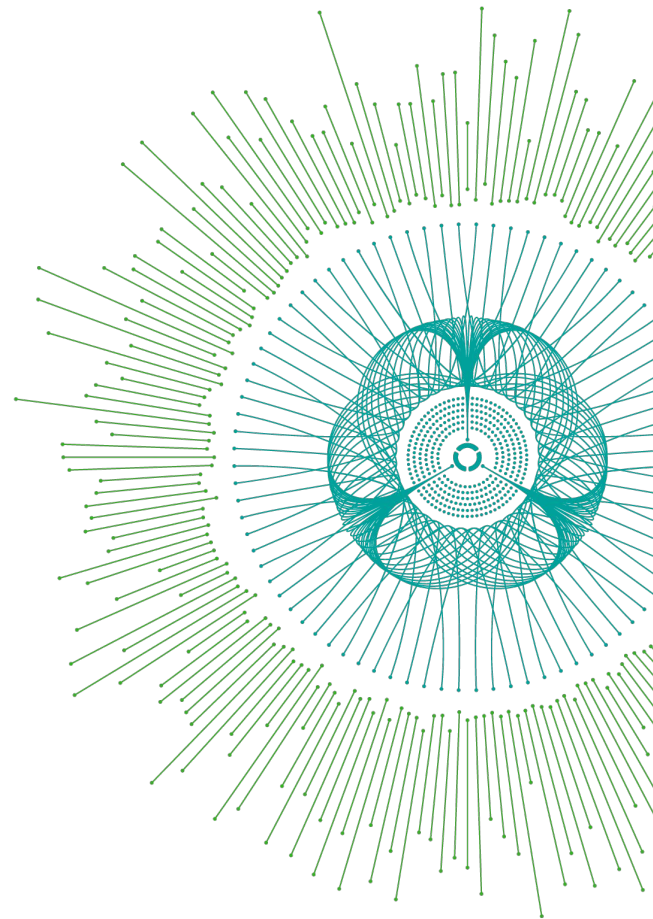




# Working with Data in the Cloud

Martin Durant

October 2020



# Contents

- 1 | Big, bigger
- 2 | Cloud storage
- 3 | Cloud-ready formats
- 4 | Compute to data
- 5 | Distributing datasets
- 6 | Summary





# Introduction: Big data, Bigger data



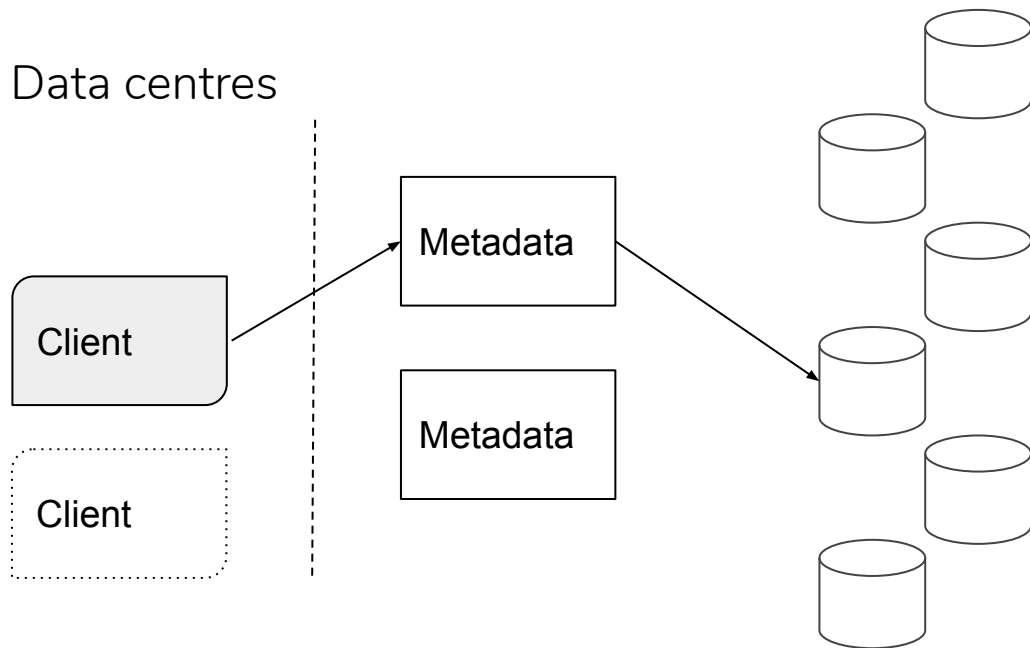
# Big, Bigger

- We are drowning in ever more data (TB->PB)
- Storing in memory is not an option
- Downloading is not an option
- Stored in variety of data farms:
  - Cloud provider object stores
  - Cluster storage (HDFS)
  - Institutional servers (NFS etc)
- Auth and access



# Big, Bigger

- Data centres



# Big, Bigger: caveats

- Object stores are not filesystems
- Every system has its own ideas about auth
- Access from within the network very different from without





fspec



# fsspec: introduction

- Consistent API over many storage backend
- Explore any store like local files
- Makes file-like objects
- Integrates with python ecosystem





# fsspec: introduction

- Implements:

- local
- http
- (s)ftp
- (web/http)HDFS
- github, jupyter, git
- AWS S3
- GCS
- MS SMB
- Azure datalake/blob
- dropbox, gdrive





# fspec: examples



# fsspec: bonus

- Convenience functions
- Path expansion
- Bulk operations
- Async/concurrency
- Caching and buffering
- Compression and text mode
- Cloud friendly





# Cloud formats



# Cloud-friendly binary formats

- A well-designed file format has the following features:
  - Human-readable metadata, with descriptive attributes
  - Binary format for efficiency and compactness
  - Choice of compression and filter operations: CPU versus size
  - Explicit data types
  - Chunks/fields of data can be accessed independently
  - **Works seamlessly with remote/cloud storage**



# Cloud-friendly?

- JSON { }
- XML < / >
- CSV , ,
  - with schema?
- Excel
- HDF
- pickle
- avro/thrift/protobufs
- proprietary/custom



# Cloud-friendly? Excel??



```
import pandas as pd
pd.read_excel(
    "zip://FinancialSample.xlsx::"
    "s3://mymdtemp/FinancialSample.xlsx.zip"
)
```



# Cloud-friendly: parquet and zarr

- Efficient binary encoding and compression
- Metadata stored separately
- Strongly typed
- Natural chunking - load what you need
- Filename conventions







# Cloud formats: examples





# Cataloging



# Cataloging

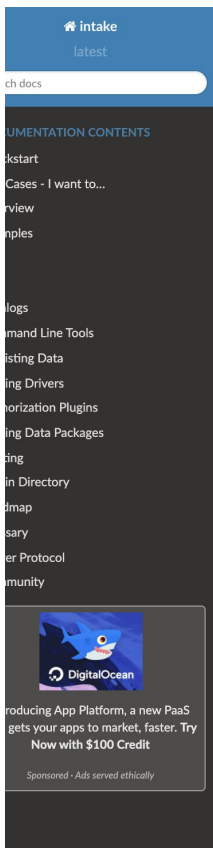


# INTAKE

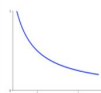
- Consistent API to all data sources
- Hierarchical, searchable catalogue trees
- Metadata lives with dataset declaration
- Intake allows data as declarative code, you can
  - Share (files, remote)
  - Update in place
  - Version control
  - Group and structure
  - Package and distribute



# Cataloging



## Data User



- Intake loads the data for a range of formats and types (see [Plugin Directory](#)) into containers you already use, like Pandas dataframes, Python lists, NumPy arrays, and more
- Intake loads, then gets out of your way
- GUI search and introspect data-sets in [Catalogs](#): quickly find what you need to do your work

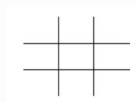
- Install data-sets and automatically get requirements
- Leverage cloud resources and distributed computing.

See the executable tutorial:



## Data Provider

- Simple spec to define data sources
- Single point of truth, no more copy&paste
- Distribute data using packages, shared files or a server
- Update definitions in-place
- Parametrise user options
- Make use of additional functionality like filename parsing and caching.



See the executable tutorial:



## IT

- Create catalogs out of established departmental practices
- Provide data access credentials via Intake parameters
- Use server-client architecture as gatekeeper:
  - add authentication methods
  - add monitoring point; track the data-sets being accessed.



- Hook Intake into proprietary data access systems.

## Developer



# INTAKE



# Making a catalogue

- Figure out how to load
- Encode the spec
- Save to file
- Add detail
- Done
- Update!



# Share catalogue

- Send the file
- Upload the file, add access or make public
  - network drive
  - github
  - dropbox/gdrive
  - cloud providers (azure, aws, gcp)
  - file server
- Build a package





# Intake Example



# Compute to data

- Code is small
- Results are (usually) small
- Easily scale up or out
- Bandwidth
- Familiar browser UI
- Sharing







# Summary



# Summary

- Big data:
  - store in the cloud or network resources
  - `fsspec` allows uniform experience for all backends
- Cloud friendly formats
  - chunking
- Metadata and cataloging
  - `Intake` to reference all your data sources
- Compute to data



Thank You!

